



Department of Computer Science and Engineering (Data Science)

Subject: Applied Data Science (DJ19DSL703)

Experiment - 6

(Model Building)

SAP ID: 60009200056

NAME: Rishabh Patil

DIV/BATCH: D12

Aim: To build model using feature engineering.

Theory:

Feature engineering is a critical step in the data pre-processing process in data science and machine learning. It involves creating new features or modifying existing ones in your dataset to improve the performance of machine learning models. Proper feature engineering can lead to more accurate and robust models. Here are some key aspects of feature engineering:

- **Feature Extraction:** Feature extraction involves transforming raw data into a format that's suitable for machine learning. For example, extracting date features (e.g., year, month, day) from a timestamp, converting text data into numerical representations (e.g., TF-IDF, word embedding), or summarizing information in image data (e.g., color histograms, edge detection).
- **Feature Transformation:** This involves applying mathematical or statistical transformations to your features to make them more suitable for modelling. Common techniques include scaling (e.g., standardization or normalization), log or power transformations, and encoding categorical variables (e.g., one-hot encoding or label encoding).
- **Feature Creation:** Sometimes, you may need to create new features based on domain knowledge or insights gained during data exploration. This could involve combining existing features, creating interaction terms, or engineering new variables to capture specific patterns or relationships.
- **Handling Missing Data:** Dealing with missing data is also part of feature engineering. You can choose to impute missing values using techniques like mean, median, or predictive modelling. Sometimes, you may create binary flags indicating the presence of missing data.
- **Feature Selection:** Feature engineering also involves selecting the most relevant features for your model. This can be done through techniques like univariate feature selection, feature importance from tree-based models, or through domain knowledge.
- **Text and NLP Feature Engineering:** When working with text data, you may need to perform additional feature engineering, such as tokenization, stemming, lemmatization, and sentiment analysis.
- **Handling Categorical Data:** Categorical variables require special attention. You can use techniques like one-hot encoding, label encoding, or target encoding to represent categorical data numerically.
- **Temporal Feature Engineering:** For time-series data, creating lag features or time-based aggregations can be valuable.
- **Geospatial Feature Engineering:** For geospatial data, you might calculate distances, create spatial clusters, or derive location-based statistics.



Department of Computer Science and Engineering (Data Science)

- **Domain-Specific Feature Engineering:** In some cases, domain-specific knowledge can lead to unique feature engineering approaches. This might include creating custom metrics or indicators specific to your problem.
- Feature engineering is an iterative process, and it often requires experimenting with different feature combinations and transformations to find the most informative and predictive features for your machine learning model. Effective feature engineering can significantly impact the model's performance and its ability to uncover valuable insights from the data.

Model Evaluation / Performance Metrics

The choice of performance metrics depends on the type of machine learning problem. Here are some common metrics for different types of problems:

- For binary classification, you can use metrics like accuracy, precision, recall, F1-score, and the receiver operating characteristic (ROC) curve.
- For multi-class classification, metrics like accuracy, precision, recall, F1-score, and confusion matrices can be useful.
- For regression problems, metrics like mean squared error (MSE), mean absolute error (MAE), and R-squared (R²) are commonly used.

Identify and Minimize data leakage

•**Train-Test Split:** This is a common practice in which you divide your dataset into two subsets: one for training the model and the other for testing its performance. The split, often referred to as the training set and the test set, allows you to train your model on a portion of the data and then evaluate its performance on unseen data.

•**Cross-Validation:** Cross-validation is a more robust technique than a simple train-test split. It involves partitioning the dataset into multiple subsets (usually k subsets or "folds") and training and evaluating the model multiple times. The results are then averaged to provide a more reliable estimate of the model's performance. Common cross-validation methods include k-fold cross-validation and stratified k-fold cross-validation.

Lab Assignment:

- At this stage, we should be having a cleaned data. We will consider the cleaned data for feature engineering steps. Identify new features, make existing features better, remove unwanted or highly correlated features.
- We will perform train test split of the data. Train the model on the training data and evaluate the model on test data.
- Perform cross validation (changing train test split data) and then rechecking metrics to see if the model is not overfitting.

Assignment Examples to work with

1. House price prediction
2. Sales prediction
3. Credit card fraud detection
4. Customer segmentation
5. Diabetes prediction

```
In [ ]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore")

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error

from xgboost import XGBRegressor
```


```
In [ ]: df = pd.read_csv("HousePricePrediction.csv")
```

```
In [ ]: df.head()
```

```
Out[ ]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utili
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	All
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	All
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	All
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	All
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	All

5 rows × 81 columns



Mapping Statistics of Data

In []: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1460 entries, 0 to 1459
```

```
Data columns (total 81 columns):
```

#	Column	Non-Null Count	Dtype
0	Id	1460 non-null	int64
1	MSSubClass	1460 non-null	int64
2	MSZoning	1460 non-null	object
3	LotFrontage	1201 non-null	float64
4	LotArea	1460 non-null	int64
5	Street	1460 non-null	object
6	Alley	91 non-null	object
7	LotShape	1460 non-null	object
8	LandContour	1460 non-null	object
9	Utilities	1460 non-null	object
10	LotConfig	1460 non-null	object
11	LandSlope	1460 non-null	object
12	Neighborhood	1460 non-null	object
13	Condition1	1460 non-null	object
14	Condition2	1460 non-null	object
15	BldgType	1460 non-null	object
16	HouseStyle	1460 non-null	object
17	OverallQual	1460 non-null	int64
18	OverallCond	1460 non-null	int64
19	YearBuilt	1460 non-null	int64
20	YearRemodAdd	1460 non-null	int64
21	RoofStyle	1460 non-null	object
22	RoofMatl	1460 non-null	object
23	Exterior1st	1460 non-null	object
24	Exterior2nd	1460 non-null	object
25	MasVnrType	1452 non-null	object
26	MasVnrArea	1452 non-null	float64
27	ExterQual	1460 non-null	object
28	ExterCond	1460 non-null	object
29	Foundation	1460 non-null	object
30	BsmtQual	1423 non-null	object
31	BsmtCond	1423 non-null	object
32	BsmtExposure	1422 non-null	object
33	BsmtFinType1	1423 non-null	object
34	BsmtFinSF1	1460 non-null	int64
35	BsmtFinType2	1422 non-null	object
36	BsmtFinSF2	1460 non-null	int64
37	BsmtUnfSF	1460 non-null	int64
38	TotalBsmtSF	1460 non-null	int64
39	Heating	1460 non-null	object
40	HeatingQC	1460 non-null	object
41	CentralAir	1460 non-null	object
42	Electrical	1459 non-null	object
43	1stFlrSF	1460 non-null	int64
44	2ndFlrSF	1460 non-null	int64
45	LowQualFinSF	1460 non-null	int64
46	GrLivArea	1460 non-null	int64
47	BsmtFullBath	1460 non-null	int64
48	BsmtHalfBath	1460 non-null	int64
49	FullBath	1460 non-null	int64
50	HalfBath	1460 non-null	int64
51	BedroomAbvGr	1460 non-null	int64

52	KitchenAbvGr	1460	non-null	int64
53	KitchenQual	1460	non-null	object
54	TotRmsAbvGrd	1460	non-null	int64
55	Functional	1460	non-null	object
56	Fireplaces	1460	non-null	int64
57	FireplaceQu	770	non-null	object
58	GarageType	1379	non-null	object
59	GarageYrBlt	1379	non-null	float64
60	GarageFinish	1379	non-null	object
61	GarageCars	1460	non-null	int64
62	GarageArea	1460	non-null	int64
63	GarageQual	1379	non-null	object
64	GarageCond	1379	non-null	object
65	PavedDrive	1460	non-null	object
66	WoodDeckSF	1460	non-null	int64
67	OpenPorchSF	1460	non-null	int64
68	EnclosedPorch	1460	non-null	int64
69	3SsnPorch	1460	non-null	int64
70	ScreenPorch	1460	non-null	int64
71	PoolArea	1460	non-null	int64
72	PoolQC	7	non-null	object
73	Fence	281	non-null	object
74	MiscFeature	54	non-null	object
75	MiscVal	1460	non-null	int64
76	MoSold	1460	non-null	int64
77	YrSold	1460	non-null	int64
78	SaleType	1460	non-null	object
79	SaleCondition	1460	non-null	object
80	SalePrice	1460	non-null	int64

dtypes: float64(3), int64(35), object(43)
memory usage: 924.0+ KB

In []: `df.describe().T`


Out[]:

	count	mean	std	min	25%	50%	75%	
Id	1460.0	730.500000	421.610009	1.0	365.75	730.5	1095.25	
MSSubClass	1460.0	56.897260	42.300571	20.0	20.00	50.0	70.00	
LotFrontage	1201.0	70.049958	24.284752	21.0	59.00	69.0	80.00	
LotArea	1460.0	10516.828082	9981.264932	1300.0	7553.50	9478.5	11601.50	21
OverallQual	1460.0	6.099315	1.382997	1.0	5.00	6.0	7.00	
OverallCond	1460.0	5.575342	1.112799	1.0	5.00	5.0	6.00	
YearBuilt	1460.0	1971.267808	30.202904	1872.0	1954.00	1973.0	2000.00	
YearRemodAdd	1460.0	1984.865753	20.645407	1950.0	1967.00	1994.0	2004.00	
MasVnrArea	1452.0	103.685262	181.066207	0.0	0.00	0.0	166.00	
BsmtFinSF1	1460.0	443.639726	456.098091	0.0	0.00	383.5	712.25	
BsmtFinSF2	1460.0	46.549315	161.319273	0.0	0.00	0.0	0.00	
BsmtUnfSF	1460.0	567.240411	441.866955	0.0	223.00	477.5	808.00	
TotalBsmtSF	1460.0	1057.429452	438.705324	0.0	795.75	991.5	1298.25	
1stFlrSF	1460.0	1162.626712	386.587738	334.0	882.00	1087.0	1391.25	
2ndFlrSF	1460.0	346.992466	436.528436	0.0	0.00	0.0	728.00	
LowQualFinSF	1460.0	5.844521	48.623081	0.0	0.00	0.0	0.00	
GrLivArea	1460.0	1515.463699	525.480383	334.0	1129.50	1464.0	1776.75	
BsmtFullBath	1460.0	0.425342	0.518911	0.0	0.00	0.0	1.00	
BsmtHalfBath	1460.0	0.057534	0.238753	0.0	0.00	0.0	0.00	
FullBath	1460.0	1.565068	0.550916	0.0	1.00	2.0	2.00	
HalfBath	1460.0	0.382877	0.502885	0.0	0.00	0.0	1.00	
BedroomAbvGr	1460.0	2.866438	0.815778	0.0	2.00	3.0	3.00	
KitchenAbvGr	1460.0	1.046575	0.220338	0.0	1.00	1.0	1.00	
TotRmsAbvGrd	1460.0	6.517808	1.625393	2.0	5.00	6.0	7.00	
Fireplaces	1460.0	0.613014	0.644666	0.0	0.00	1.0	1.00	
GarageYrBlt	1379.0	1978.506164	24.689725	1900.0	1961.00	1980.0	2002.00	
GarageCars	1460.0	1.767123	0.747315	0.0	1.00	2.0	2.00	
GarageArea	1460.0	472.980137	213.804841	0.0	334.50	480.0	576.00	
WoodDeckSF	1460.0	94.244521	125.338794	0.0	0.00	0.0	168.00	
OpenPorchSF	1460.0	46.660274	66.256028	0.0	0.00	25.0	68.00	
EnclosedPorch	1460.0	21.954110	61.119149	0.0	0.00	0.0	0.00	
3SsnPorch	1460.0	3.409589	29.317331	0.0	0.00	0.0	0.00	
ScreenPorch	1460.0	15.060959	55.757415	0.0	0.00	0.0	0.00	
PoolArea	1460.0	2.758904	40.177307	0.0	0.00	0.0	0.00	
MiscVal	1460.0	43.489041	496.123024	0.0	0.00	0.0	0.00	1
MoSold	1460.0	6.321918	2.703626	1.0	5.00	6.0	8.00	

	count	mean	std	min	25%	50%	75%
YrSold	1460.0	2007.815753	1.328095	2006.0	2007.00	2008.0	2009.00
SalePrice	1460.0	180921.195890	79442.502883	34900.0	129975.00	163000.0	214000.00

Correlation Features

```
In [ ]: plt.figure(figsize=(10,8))
sns.heatmap(df.corr(), cmap="RdBu")
plt.title("Correlations Between Variables", size=15)
plt.show()
```




Feature Selection

```
In [ ]: important_num_cols = list(df.corr()["SalePrice"][(df.corr()["SalePrice"]>0.50)
| (df.corr()["SalePrice"]<-0.50)].index)
# Only those columns are selected which can be important (Subjective)
cat_cols = ["MSZoning", "Utilities", "BldgType", "Heating", "KitchenQual", "SaleCo
ndition", "LandSlope"]
important_cols = important_num_cols + cat_cols
df = df[important_cols]
df
```

Out[]:

	OverallQual	YearBuilt	YearRemodAdd	TotalBsmtSF	1stFlrSF	GrLivArea	FullBath	TotRms
0	7	2003	2003	856	856	1710	2	
1	6	1976	1976	1262	1262	1262	2	
2	7	2001	2002	920	920	1786	2	
3	7	1915	1970	756	961	1717	1	
4	8	2000	2000	1145	1145	2198	2	
...
1455	6	1999	2000	953	953	1647	2	
1456	6	1978	1988	1542	2073	2073	2	
1457	7	1941	2006	1152	1188	2340	2	
1458	5	1950	1996	1078	1078	1078	1	
1459	5	1965	1965	1256	1256	1256	1	

1460 rows × 18 columns



```
In [ ]: print("Missing Values by Column")
print("-"*30)
print(df.isna().sum())
print("-"*30)
print("TOTAL MISSING VALUES:",df.isna().sum().sum())
```

Missing Values by Column


```
-----
OverallQual      0
YearBuilt        0
YearRemodAdd     0
TotalBsmtSF      0
1stFlrSF         0
GrLivArea        0
FullBath         0
TotRmsAbvGrd    0
GarageCars       0
GarageArea       0
SalePrice        0
MSZoning         0
Utilities        0
BldgType         0
Heating          0
KitchenQual      0
SaleCondition    0
LandSlope        0
dtype: int64
```


```
-----
TOTAL MISSING VALUES: 0
```


Correlation Using Joint Plots


```
In [ ]: plt.figure(figsize=(10,8))
        for i in df.columns:
            if i not in cat_cols:
                sns.jointplot(x=df[i], y=df["SalePrice"])
                plt.show()
```


<Figure size 1000x800 with 0 Axes>
















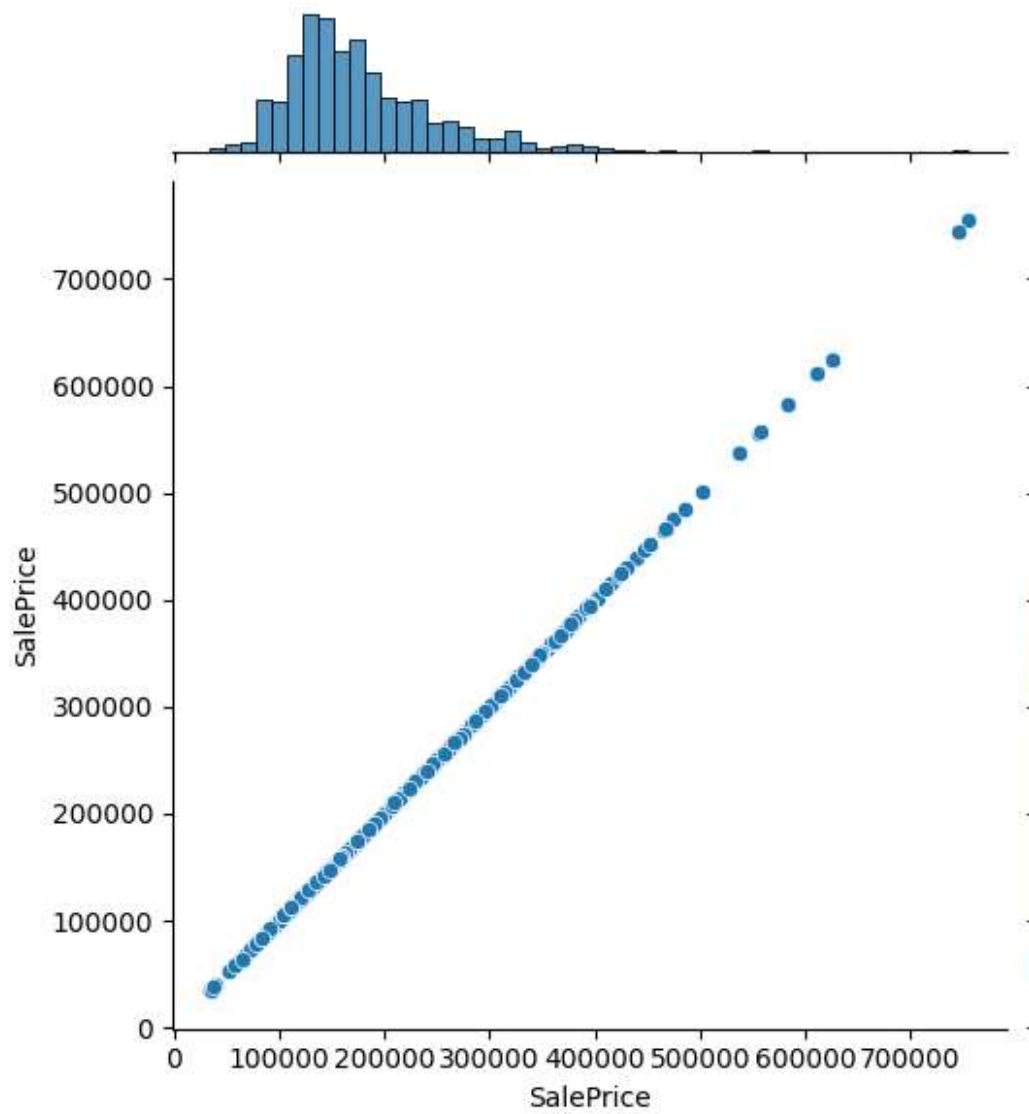













One hot encoding

```
In [ ]: df = pd.get_dummies(df, columns=cat_cols)
df
```

Out[]:

	OverallQual	YearBuilt	YearRemodAdd	TotalBsmtSF	1stFlrSF	GrLivArea	FullBath	TotRms
0	7	2003	2003	856	856	1710	2	
1	6	1976	1976	1262	1262	1262	2	
2	7	2001	2002	920	920	1786	2	
3	7	1915	1970	756	961	1717	1	
4	8	2000	2000	1145	1145	2198	2	
...
1455	6	1999	2000	953	953	1647	2	
1456	6	1978	1988	1542	2073	2073	2	
1457	7	1941	2006	1152	1188	2340	2	
1458	5	1950	1996	1078	1078	1078	1	
1459	5	1965	1965	1256	1256	1256	1	

1460 rows × 42 columns



Scaling the numerical values



```
In [ ]: important_num_cols.remove("SalePrice")

scaler = StandardScaler()
df[important_num_cols] = scaler.fit_transform(df[important_num_cols])
df
```

Out[]:

	OverallQual	YearBuilt	YearRemodAdd	TotalBsmtSF	1stFlrSF	GrLivArea	FullBath	TotR
0	0.651479	1.050994	0.878668	-0.459303	-0.793434	0.370333	0.789741	
1	-0.071836	0.156734	-0.429577	0.466465	0.257140	-0.482512	0.789741	
2	0.651479	0.984752	0.830215	-0.313369	-0.627826	0.515013	0.789741	
3	0.651479	-1.863632	-0.720298	-0.687324	-0.521734	0.383659	-1.026041	
4	1.374795	0.951632	0.733308	0.199680	-0.045611	1.299326	0.789741	
...	
1455	-0.071836	0.918511	0.733308	-0.238122	-0.542435	0.250402	0.789741	
1456	-0.071836	0.222975	0.151865	1.104925	2.355701	1.061367	0.789741	
1457	0.651479	-1.002492	1.024029	0.215641	0.065656	1.569647	0.789741	
1458	-0.795151	-0.704406	0.539493	0.046905	-0.218982	-0.832788	-1.026041	
1459	-0.795151	-0.207594	-0.962566	0.452784	0.241615	-0.493934	-1.026041	

1460 rows × 42 columns



Train Test Split

```
In [ ]: X = df.drop("SalePrice", axis=1)
        y = df["SalePrice"]
```

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Useful Metrics Calculations

```
In [ ]: def evaluation(y, predictions):
        mae = mean_absolute_error(y, predictions)
        mse = mean_squared_error(y, predictions)
        rmse = np.sqrt(mean_squared_error(y, predictions))
        r_squared = r2_score(y, predictions)
        return mae, mse, rmse, r_squared
```

Modelling XGBoostRegressor

```
In [ ]: xgb = XGBRegressor(n_estimators=1000, learning_rate=0.01)
xgb.fit(X_train, y_train)
predictions = xgb.predict(X_test)

mae, mse, rmse, r_squared = evaluation(y_test, predictions)
print("MAE:", mae)
print("MSE:", mse)
print("RMSE:", rmse)
print("R2 Score:", r_squared)
```

MAE: 17985.48701038099
MSE: 837038411.4275047
RMSE: 28931.616121943563
R2 Score: 0.8908731664315402

Cross-Validation

```
In [ ]: test_score = xgb.score(X_test, y_test)
test_score
```

Out[]: 0.8908731664315402

```
In [ ]: k = 5
cross_val_scores = cross_val_score(xgb, X_train, y_train, cv=k)

print("Cross-Validation Scores:", cross_val_scores)
print("Mean CV Score:", np.mean(cross_val_scores))
```

Cross-Validation Scores: [0.82876146 0.67764145 0.86030628 0.85252966 0.86836063]
Mean CV Score: 0.8175198945132556

```
In [ ]: if np.mean(cross_val_scores) > test_score:
        print("The model may be overfitting.")
else:
        print("The model is not overfitting.")
```

The model is not overfitting.

```
In [ ]:
```